

Interop

bridging the divide between managed and native code

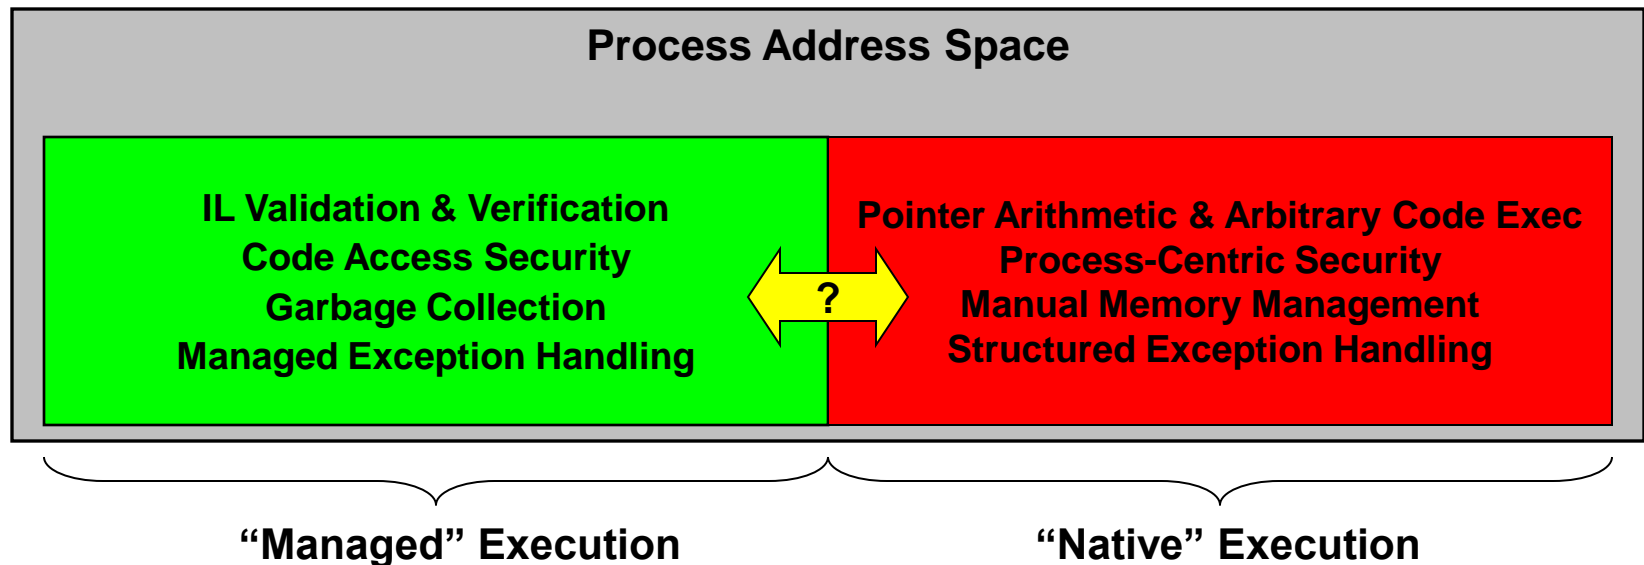


Outline

- **CLR support for managed/native code interop**
 - Comparing “managed execution” and “native execution”
 - The metadata-driven partnership
 - Mechanics of managed/native interop
 - CLR => Win32
 - CLR => COM
 - COM => CLR

Comparing “Managed” & “Native” Execution

- In practice, .NET apps involve a mix of managed & native code
 - CLR-based “managed code”
 - Win32/COM-based “native code”
 - Transitions between the two worlds must be carefully coordinated
 - it’s more than simply parameter marshaling



The Metadata-Driven Partnership

- **Metadata (type information) drives managed/native interop**
 - Managed type info can be derived from native type info
 - Native type info can be derived from managed type info
 - Pro: tools can automate **most** (not all) transformations
 - Con: tools can automate most (**not all**) transformations
- **Interop is a partnership involving you, the compiler, and the CLR**
 - The CLR (and its tools) will automate as much as possible
 - limited only by the fidelity of the native type information that's available
 - You may have to fill in some blanks or fine tune some transformations

Interop Facilities


- **The CLR supports two kinds of managed/native interop**
 - Interop with Win32 DLLs
 - Platform Invocation Services
 - "P/Invoke"
 - Interop with COM components
 - "COM interop"

P/Invoke

- **Managed code can load & call into Win32 DLLs**
 - Type information for Win32 DLLs is sorely lacking
 - export table contains only names & relative addresses of symbols
 - does not indicate whether exported symbol is a function or global variable
 - does not indicate “shape” of functions (its signature) or variables (its type)
 - The partnership
 - programmer describes function location & signature in terms of CLR types
 - managed compiler produces assembly/type metadata that drives JIT compilation
 - JIT compiler uses metadata to build “stubs” that perform CLR/Win32 transition
 - load the requested DLL (LoadLibrary)
 - locate the target function in memory (GetProcAddress)
 - marshal parameters to/from the target function

P/Invoke Mechanics

- The CLR/Win32 transition is carried out by *stubs* & *helpers*
 - *stubs* are little chunks of native code emitted by the JIT compiler
 - specific to the marshaling requirements dictated by the p/invoke declaration
 - *helpers* are native functions typically found in mscor*.dll
 - general purpose functions that don't need to be tuned specifically to the target
 - examples:
 - convert a CLR System.String parameter into NUL-terminated ANSI string
 - given a collection of CAS permissions, perform a demand/assert/etc



P/Invoke Metadata

- **Programmer-supplied metadata drives P/Invoke**
 - Method prototype describes the native function in terms of CLR types
 - managed compiler emits metadata for the managed method (empty body)
 - managed method can be called like any other static method
 - JIT compiler uses metadata to load DLL and locate the exported method
 - JIT compiler uses metadata to build stubs that handle the transition

```
using System.Runtime.InteropServices;

class Program
{
    static void Main()
    {
        int sum = Add(2, 2);
    }

    [DllImport("nativecalc.dll")]
    static extern int Add(int a, int b);
}
```


P/Invoke Fine Tuning

- **Stub generation can be fine-tuned as needed**
 - Using properties of DllImportAttribute
 - EntryPoint
 - CharSet
 - SetLastError
 - Others
 - By applying additional attributes to parameters and/or types
 - MarshalAsAttribute
 - StructLayoutAttribute

Example: P/Invoke Fine Tuning

THE GOAL

Call this function from C#:

```
// Exported by weirdtextutils.dll
//
BSTR WINAPI Concat( wchar_t *s1, char *s2 );
```

```
using System;
using System.Runtime.InteropServices;

class Program
{
    static void Main()
    {
        string s = Concat("Hello, ", "world!");
        Console.WriteLine(s);
    }

    [DllImport("weirdtextutils.dll")]
    [return: MarshalAs(UnmanagedType.BStr)]
    static extern string Concat(
        [MarshalAs(UnmanagedType.LPWSTR)] string s1,
        [MarshalAs(UnmanagedType.LPStr)] string s2
    );
}
```

Marshal the
return value
as a Basic string

Marshal s1 as a
Unicode string.


Marshal s2 as an
ANSI string.

COM Interop: CLR=>COM

- **Managed code can interact fairly easily with COM code**
 - COM type information (TLB) is fairly complete
- **The partnership**
 - TLBIMP.EXE translates COM type info into CLR type info
 - .TLB => TLBIMP.EXE => .DLL ("interop assembly")
 - COM coclass => CLR class (concrete)
 - COM interface => CLR interface
 - automated by VS.NET Add Reference wizard
 - Programmer adds a reference to interop assembly
 - CLR/JIT compiler use metadata to construct "runtime-callable wrappers"
 - RCWs

Runtime-Callable Wrappers (RCW)

- RCWs bridge the divide between two very different worlds



Instantiation: `operator new`

CLSID/ProgId registry mappings
`CoCreateInstance`

Type navigation: `is/as/cast`

`IUnknown::QueryInterface`

Error handling: `System.Exception`
`throw/try/catch/finally`

`IErrorInfo`
`(Set|Get)ErrorInfo`
Win32 SEH

Memory mgmt: `shared heap`
`garbage collection`

`CoTaskMem(Alloc|Free)`
`IUnknown::AddRef/Release`

COM Type Information Deficiencies

- COM type information (TLB) does not quite offer full fidelity
- Example
 - COM IDL supports C-style “conformant arrays”
 - one parameter is a pointer to the first element in an array
 - another parameter specifies the element count
 - COM type libraries cannot represent conformant arrays
 - result is a pointer to a single entity, not the start of an array of entities

```
interface ICalc : IUnknown
{
    HRESULT Sum( [in, size_is(count)] long *pValues,
                [in] long count,
                [out, retval] long *pResult );
};
```

IDL

midl.exe

```
interface ICalc : IUnknown
{
    HRESULT Sum( [in] long *pValues,
                [in] long count,
                [out, retval] long *pResult );
};
```

TLB

tlbimp.exe

interop
assembly

```
public interface ICalc
{
    int Sum( ref int pValues,
            int count );
};
```

Wrong!

COM Interop Fine Tuning

- RCW metadata can be provided manually when needed ala P/Invoke
 - programmer manually describes types in CLR terms
 - attributes are used to provide required instantiation and marshaling details
 - CLR/JIT compiler uses programmer-supplied metadata to construct RCW

```
[ComImport, Guid("37DE3F74-99BE-4DD9-A06D-422203752987") ]
```

```
class CalcClass
```

```
{  
    // empty  
}
```

Value taken from IDL/TLB

```
[ Guid("22E8E9BF-552E-4A09-8B93-33778020F240"),  
  InterfaceType(ComInterfaceType.InterfaceIsIUnknown) ]
```

```
public interface ICalc
```

```
{  
    int Sum(  
        [MarshalAs(UnmanagedType.LPArray, SizeParamIndex = 1)] int[] values,  
        int count  
    );  
}
```

Adjust how RCW marshals the "values" parameter

```
class Program
```

```
{  
    static void Main()  
    {  
        ICalc c = (ICalc)new CalcClass();  
        int[] values = { 1, 2, 3, 4 };  
        int sum = c.Sum(values, values.Length);  
    }  
}
```

COM & Threads

- **COM has an “apartment model” construct for dealing with threads**
 - Components declare their preparedness/requirements re calling threads
 - ThreadingModel registry setting
 - Calling threads declare their preparedness/requirements re components
 - CoInitialize(Ex)
 - COM activation returns proxies or not based on those two settings
 - CoCreateInstance(Ex)
- **CLR threads have a property that declares their COM apartment state**
 - System.Threading.Thread.ApartmentState
 - System.STAThreadAttribute
 - System.MTAThreadAttribute

```
using System;


class Program
{
    [MTAThread]
    static void Main()
    {
        ICalc c = (ICalc)new CalcClass();
        int sum = c.Add(2, 3);
    }
}
```

COM Interop: COM => CLR


- **CLR type information offers high fidelity**
 - CLR-to-COM metadata transformations are often more easily performed
 - Does not mean every valid type can be represented in COM (e.g.: generics)
- **The partnership**
 - TLBEXP.EXE can be used to translate CLR type info into COM type info
 - .DLL (assembly) => TLBEXP.EXE => .TLB
 - suitable for use by VB6, MSFT C++ & other TLB-aware compilers
 - REGASM.EXE performs COM-required registration
 - all roads lead to MSCOREE.DLL
 - MSCOREE.DLL builds COM-callable wrappers (CCWs)
 - driven by metadata in .NET assembly (not the TLB)
 - .NET programmer can use attributes to fine tune tool operation

COM-Callable Wrappers (CCW)

- **CCWs are essentially the inverse of RCWs**
 - Constructed by mscoree.dll during COM “activation”
 - Constructed whenever CLR object references are marshaled as parameters
 - COM-induced complexity made relatively seamless by the CLR



COM Activation



VBA/VBA

```
Dim calc as Object
Set calc = CreateObject("Pluralsight.Calc")
```

```
var calc = new ActiveXObject("Pluralsight.Calc");
```

JS

Registry/HKCR/Pluralsight.Calc/CLSID

(Default) = {clsid}

ole32.dll

```
IUnknown *CoCreateInstanceEx( clsid, iid )
{
    dll = LoadLibrary(reg/Default);
    gco = GetProcAddress(dll, "DllGetClassObject");
    return gco(clsid, iid);
}
```

```
ICalc *pCalc;
CoCreateInstance( CLSID_PluralsightCalc,
                  0, CLSCTX_INPROC_SERVER,
                  IID_ICalc, (void **)&pCalc );
```

C++

mscorlib.dll

```
IUnknown *DllGetClassObject( clsid, iid )
{
    clr = StartCLR(reg/RuntimeVersion);
    asm = clr.LoadAssembly(reg/Assembly);
    type = asm.GetType(reg/Class);
    obj = asm.CreateInstance(type);
    return BuildRCWForCLRObject(type, obj);
}
```

Registry/HKCR/{clsid}/InprocServer32

```
(Default) = mscorlib.dll
ThreadingModel = Both
RuntimeVersion = v2.0.50727
Assembly = pscalc, Version=1.0.0.0, Culture=neutral, ...
Class = Pluralsight.Calc
```

COM Activation and Assembly Resolution

- **COM activation & Win32 DLL loading rules affect COM=>CLR interop**
 - MSCOREE.DLL!DllGetClassObject can be called in arbitrary process contexts
 - MSCOREE.DLL resides in %SystemRoot%\System32
 - MSCOREE.DLL needs to locate (on disk or in memory) and load assemblies
 - i.e., perform assembly resolution & everything that entails
- **Implications**
 - Four-part assembly names are stored in HKCR/{clsid}/InprocServer32
 - Assemblies intended to be used from COM should be signed & in GAC
 - Assemblies referenced/used by “top level” assembly should also be in GAC
 - REGASM.EXE can add a CodeBase hint to the registry
 - intended only for quick & dirty testing
 - always results in a warning
 - may still experience failures at runtime

Native Resource Management

- **Native resource reclamation is manual/explicit**
 - free (C), delete (C++), CloseHandle (Win32), IUnknown::Release (COM)
- **Managed resource reclamation is more intelligent**
 - Garbage collection takes care of reclaiming unreachable object memory
- **Managed classes that interact directly with native code need to help**
 - P/Invoke clients **should** override System.Object.Finalize
 - i.e., make your type *finalizable* (language-specific syntax)
 - use P/Invoke to release native resource(s)
 - do **not** interact with any managed objects you may be referencing
 - RCW & P/Invoke clients **may** implement System.IDisposable
 - enables aggressive/early native resource cleanup
 - many languages provide an exception-aware IDisposable idiom

Native Resource Management: CLR => COM

- RCWs handle the direct interaction with COM objects
 - Hold the actual interface pointer
 - Are finalizable
 - Will perform the final `IUnknown::Release` when finalized during GC
- Classes that interact with RCWs may optionally spur final Release
 - `System.Runtime.InteropServices.Marshal.ReleaseComObject`

Reference to RCW

Reference to
managed object
that implements
`IDisposable`

```
class Widget : IDisposable
{
    INativeWidget _w = new NativeWidgetClass();
    Stream _s = File.OpenWrite(@"c:\widget.log");
    bool _disposed = false;

    public void Dispose()
    {
        if (!_disposed)
        {
            Marshal.ReleaseComObject(_w);
            _s.Dispose();
            _disposed = true;
        }
    }
}
```

NOTE

`ReleaseComObject` is not just *a call* to `IUnknown::Release`, it is *the final call* to `IUnknown::Release`.

Native Resource Management: CLR => Win32

```
partial class Widget : IDisposable
{
    [DllImport("kernel32.dll")]
    static extern IntPtr CreateFileMapping(string filename);

    [DllImport("kernel32.dll")]
    static extern bool CloseHandle(IntPtr h);
}
```

p/invoke declarations
simplified for brevity

Holds a Win32 HANDLE

Reference to managed object
that implements IDisposable

release native resource
forward call to Dispose
suppresses finalization

Finalizer releases only
the native resource

```
partial class Widget : IDisposable
{
    IntPtr _h = CreateFileMapping(@"c:\widgetdata.bin");
    Stream _s = File.OpenWrite(@"c:\widget.log");
    bool _disposed = false;

    public void Dispose()
    {
        if (!_disposed)
        {
            CloseHandle(_h);
            _s.Dispose();
            GC.SuppressFinalize(this);
            _disposed = true;
        }
    }

    ~Widget()
    {
        CloseHandle(_h);
    }
}
```

Summary

- **Managed/native interop is a metadata-driven process**
 - Programmers supply all metadata for Win32 interop
 - Tools supply most metadata for COM interop
 - Attributes can be used to fine-tune stub, RCW and CCW construction
 - Some awareness of COM idioms is required
 - Leverage finalization & IDisposable as needed in first-level interop classes

References

■ Useful references

- *.NET and COM: The Complete Interoperability Guide*
 - Adam Nathan, Microsoft Principal Developer
- P/Invoke Interop Wiki
 - <http://www.pinvoke.net>
- CloudBerry S3 Utility
 - <http://www.cloudberrylab.com>