

Multithreading

introduction to concurrent execution in .NET



Overview


■ Introduction to multithreading in .NET

- Vocabulary: process versus thread
- Multi-threading use-cases
- Multi-threading caveats
- Working with threads
 - starting
 - argument passing
 - shutdown coordination
- Working with the thread pool
 - directly
 - asynchronous delegate invocation
 - asynchronous I/O


Process Anatomy

- **A process is an inert container**
 - Defines a virtual address space
 - contents not addressable from another process
 - Libraries of code are mapped into the address space
 - 1 EXE + N DLLs (dynamically loaded libraries)


Process Address Space



Process Address Space



Process Address Space




Processes & Threads

- **Threads execute code**

- A path of execution through any/all code within a single process
- Have access to any/all data within that process
 - for managed threads, within an AppDomain
- Each thread has its own callstack & copy of the CPU registers
 - technically, a CLI implementation-specific detail
- A process with no threads exits (because it can no longer perform work)

Process Address Space




Multi-threading Use-Cases

- **The benefits of leveraging multi-threading include**
 - Opportunity to scale by parallelizing CPU-bound operations
 - assuming multi-core/multi-processor hardware
 - Perform CPU-bound work while I/O operations are waiting
 - Maintain a responsive user interface
 - farming off lengthy and/or blocking operations to a separate thread
 - using thread priorities to ensure the “UI thread” has priority

Multi-threading Caveats

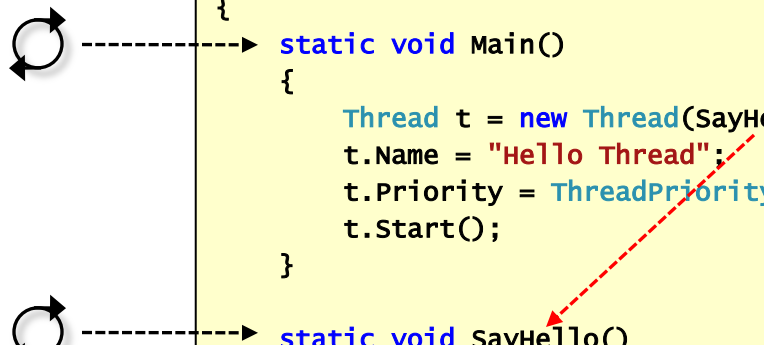
- The price to pay for multi-threading includes
 - Slower execution time on single-core/processor machines
 - context-switching overhead means...



- Added program complexity
 - lines of code
 - readability/maintainability
 - debuggability
 - testability

Starting Threads

- Threads are started explicitly using `System.Threading.Thread`
 - Constructor is used to set *thread entry point*
 - Properties can be set post-construction; prior to thread starting
 - Thread execution begins when **Start** is called



```
using System;
using System.Threading;

class Program
{
    static void Main()
    {
        Thread t = new Thread(SayHello);
        t.Name = "Hello Thread";
        t.Priority = ThreadPriority.BelowNormal;
        t.Start();
    }


    static void SayHello()
    {
        Console.WriteLine("Hello, world!");
    }
}
```

Thread Entry Point Methods

- Thread entry point methods must have one of two signatures
 - `void EntryPointMethod()`
 - MSDN/***ThreadStart***
 - `void EntryPointMethod(object stateArg)`
 - MSDN/***ParameterizedThreadStart***
 - method may be instance or static

ParameterizedThreadStart

- Thread.Start can pass an arbitrary object reference to a thread
 - “sender” and “recipient” must agree on the type of the referenced object



```
using System;
using System.Threading;

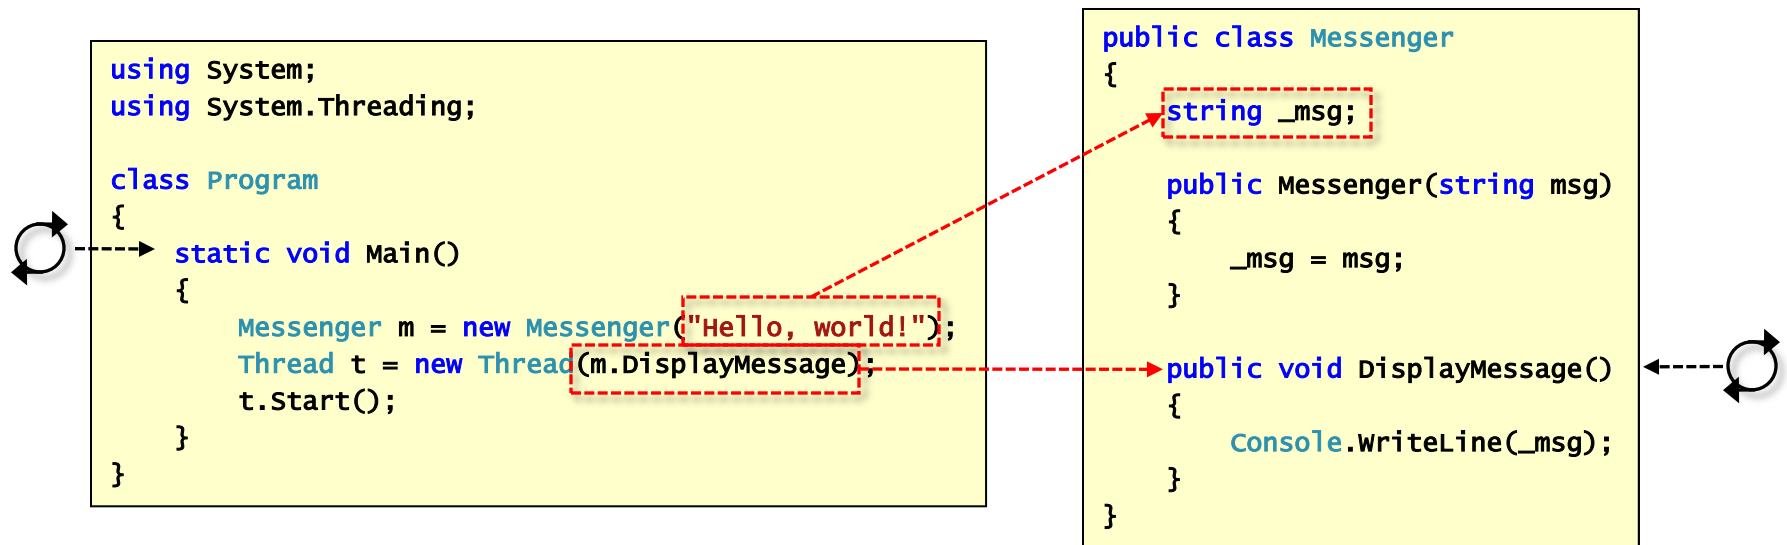
class Program
{
    static void Main()
    {
        Thread t = new Thread(DisplayMessage);
        t.Start("Hello, world!");
    }

    static void DisplayMessage(object stateArg)
    {
        string msg = stateArg as string;

        if (msg != null)
        {
            Console.WriteLine(msg);
        }
    }
}
```

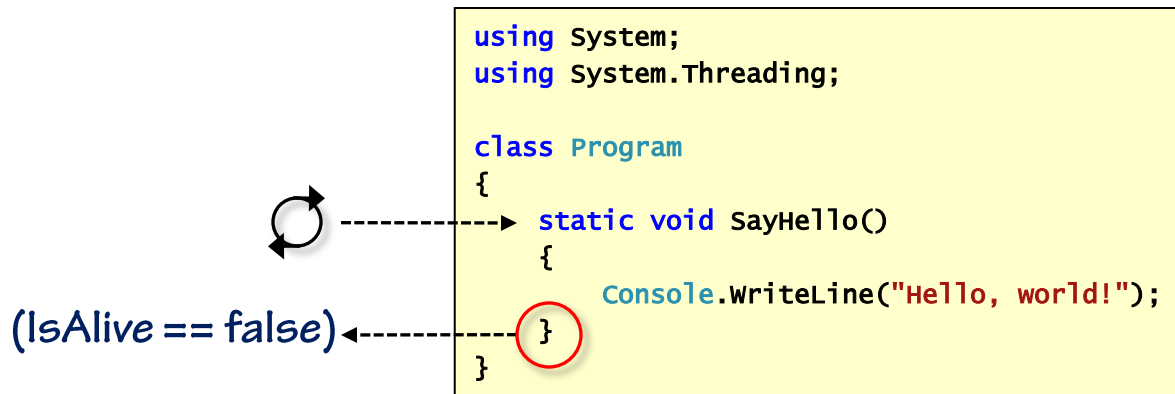
Instance Methods

- Instance methods are also suitable for thread entry point methods
 - Method has access to object fields like any other instance method would



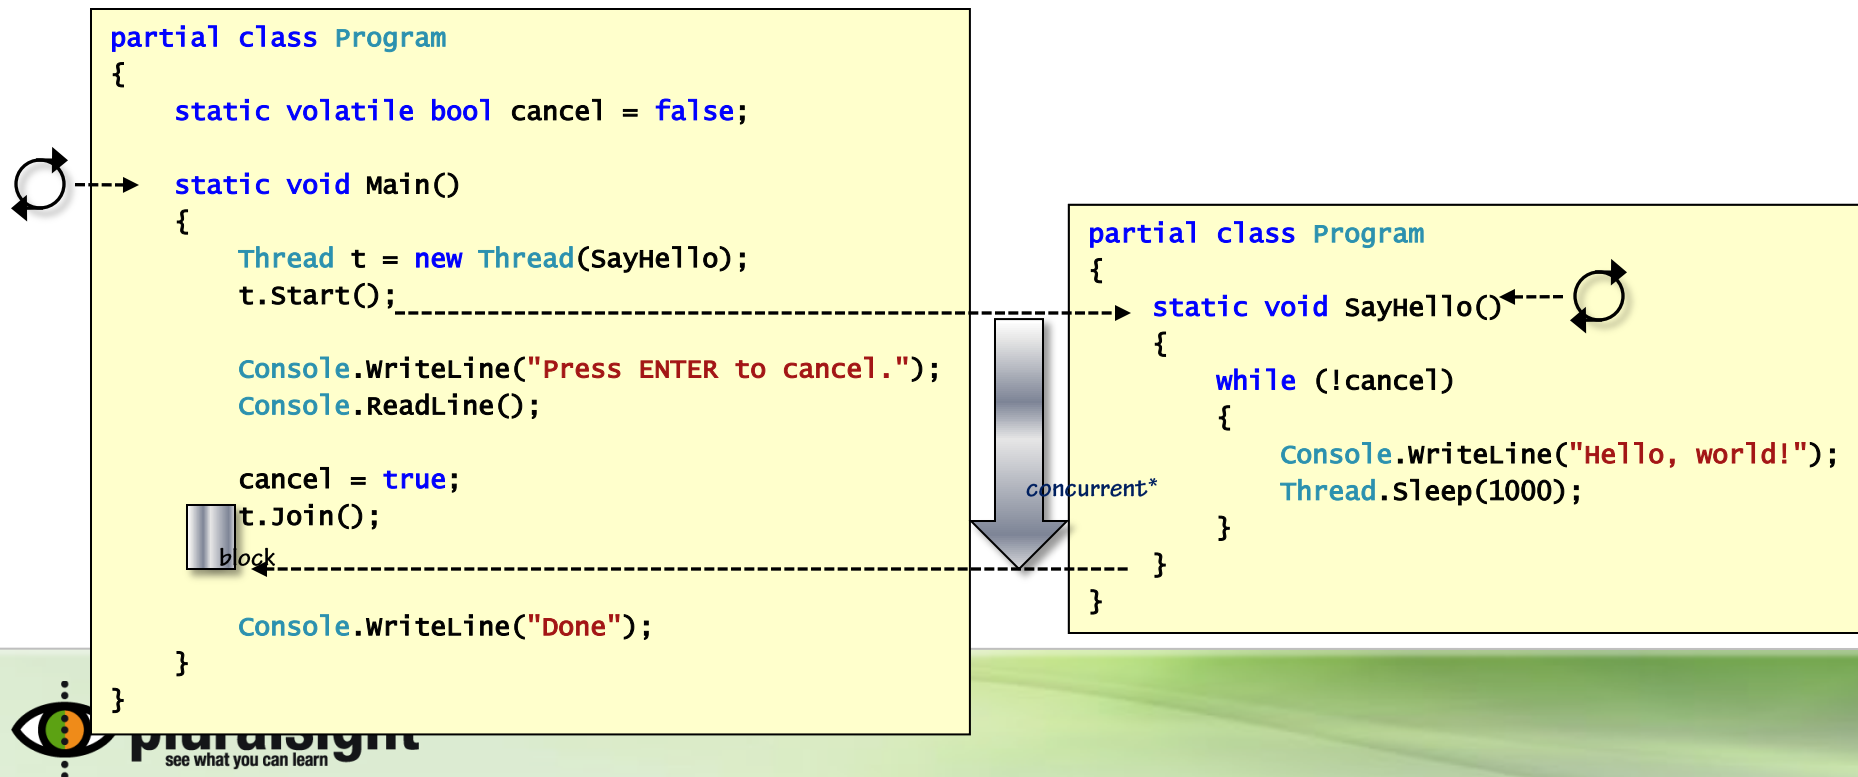
Thread Lifetime

- Execution continues until thread returns from its entry point method
 - As a result of a standard method return
 - As result of an unhandled exception
 - encountered by the thread itself ("synchronous exception")
 - induced by another thread using **Interrupt** or **Abort** ("asynchronous exception")
 - **IsAlive** provides an *instantaneous snapshot* of thread execution state



Coordinating Thread Shutdown

- Ideally, thread shutdown is choreographed
 - User-defined mechanism used to request orderly shutdown
 - Requesting thread waits until the CLR confirms the thread has exited
 - not so good: polling *IsAlive*
 - better: calling *Join*



Thread Pool

- The CLR also provides a per-process thread pool
 - Allows threads to be “borrowed” for relatively brief concurrent operations
 - CLR adds threads to, removes threads from the pool based on demand
 - Allows cost of thread start up & teardown to be amortized over life of app
 - Pooled threads have ***IsBackground*** property set to true
- Three styles of interaction with the thread pool are supported
 - ***ThreadPool.QueueUserWorkItem***
 - ***Delegate.BeginInvoke***
 - Asynchronous I/O

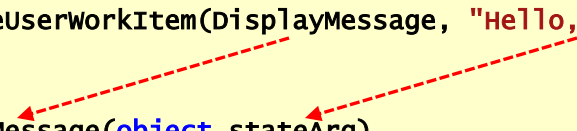
ThreadPool.QueueUserWorkItem

- Queues a request for pooled thread to call a given callback
 - Optional state argument may be passed to callback with request
 - Callback requests (and optional argument) are stored in a FIFO queue
 - Multiple reader threads means callback invocation order is **not** guaranteed

```
using System;
using System.Threading;

class Program
{
    static void Main()
    {
        ThreadPool.QueueUserWorkItem(DisplayMessage, "Hello, world!");
    }

    static void DisplayMessage(object stateArg)
    {
        Console.WriteLine(stateArg);
    }
}
```



Delegates & Async I/O

- **Delegates provide a type-aware interface to the thread pool**
 - Classes that represent a method call (signature & target instance)
 - Compiler-generated ***BeginInvoke*** method matches signature
 - CLR implements ***BeginInvoke*** at run-time to queue request to thread pool
- **Asynchronous I/O is a scalable I/O-centric interface to the thread pool**
 - QueueUserWorkItem/Delegates
 - dispatch a pooled thread to call a method that may take a long time to complete
 - Async I/O
 - queue a non-blocking I/O request to a device (***BeginRead***, ***BeginWrite***, et. al.)
 - dispatch a pooled thread to call a method notifying you *when I/O completes*

Summary

- **Threads are the core concurrent programming construct in .NET**
 - ***System.Threading.Thread***
 - Contained within/confined to a single process
 - Share access to any/all data within that process/AppDomain
 - Independently prioritized & scheduled for CPU time by the OS
 - when backed by a Win32 thread
 - Multiple techniques provided for leveraging threads
 - manual/explicit
 - thread pool
 - asynchronous I/O