

Thread Synchronization

introduction to thread synchronization in .NET



Overview

- **Introduction to thread synchronization in .NET**
 - Motivation
 - Techniques & primitives
 - atomic updates of machine word-sized values
 - data partitioning
 - wait-based synchronization primitives
 - Deadlock

Motivation

- Most resources are not meant to be accessed concurrently
 - Collections (arrays, linked-lists, etc.)
 - Files
 - ...
 - Even integers

```
partial class Program
{
    static void AddOne()
    {
        sum++;
    }
}
```



```
partial class Program
{
    static int sum = 0;

    static void Main()
    {
        Thread[] threads = new Thread[10];

        for (int n = 0; n < threads.Length; n++)
        {
            threads[n] = new Thread(AddOne);
            threads[n].Start();
        }

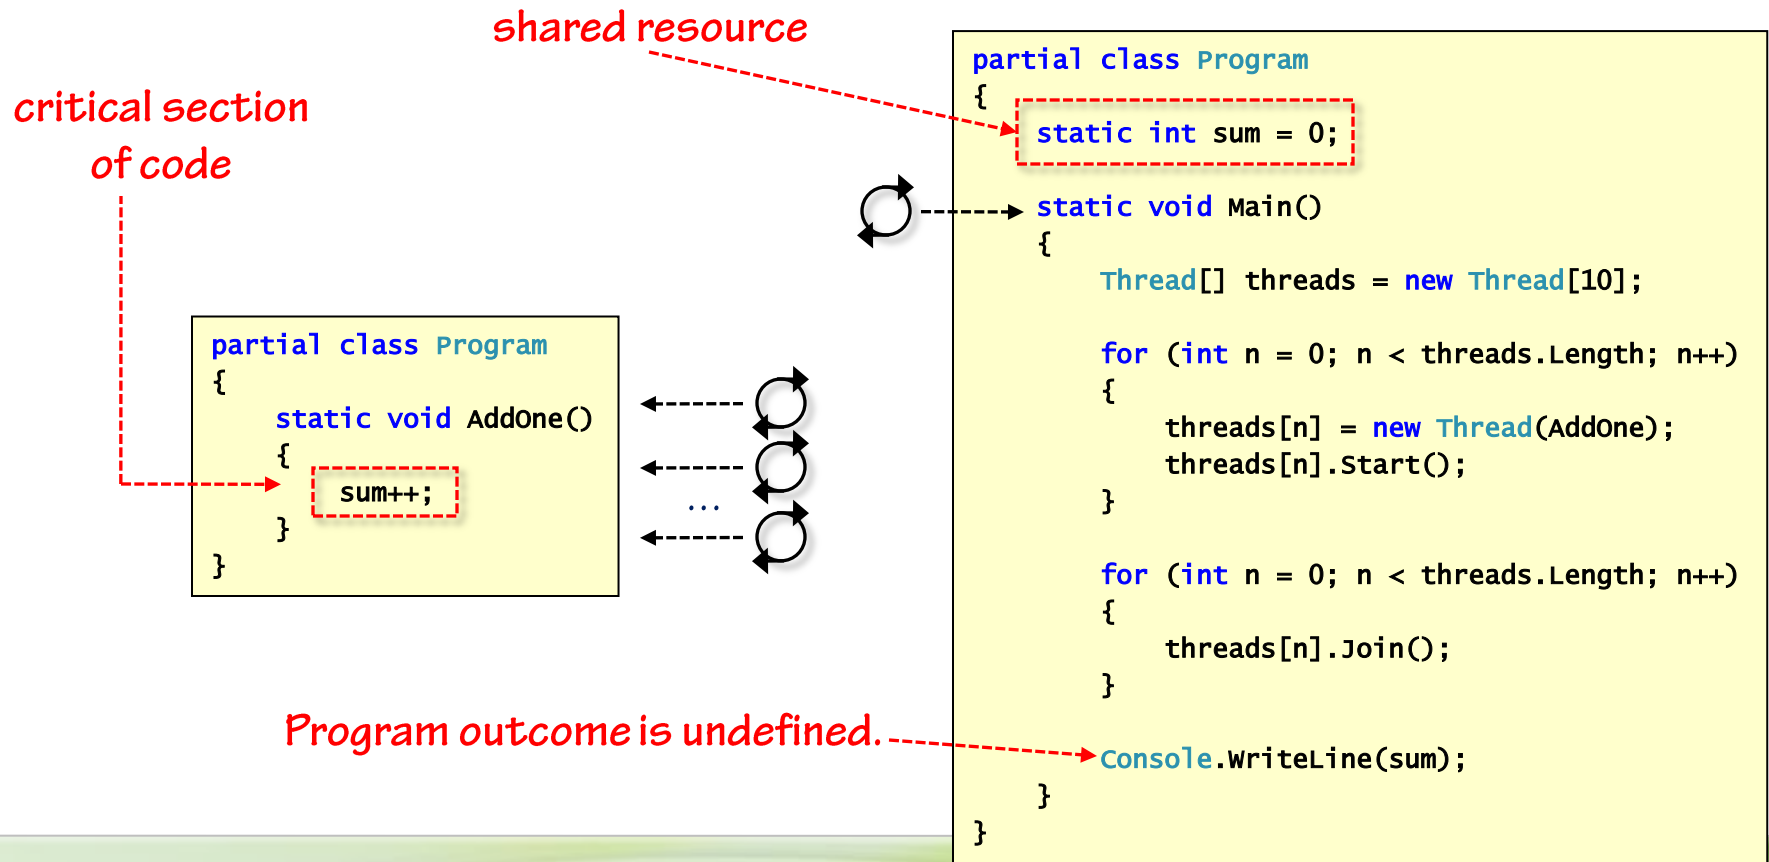
        for (int n = 0; n < threads.Length; n++)
        {
            threads[n].Join();
        }

        Console.WriteLine(sum);
    }
}
```

What should this program display?
What will this program display?

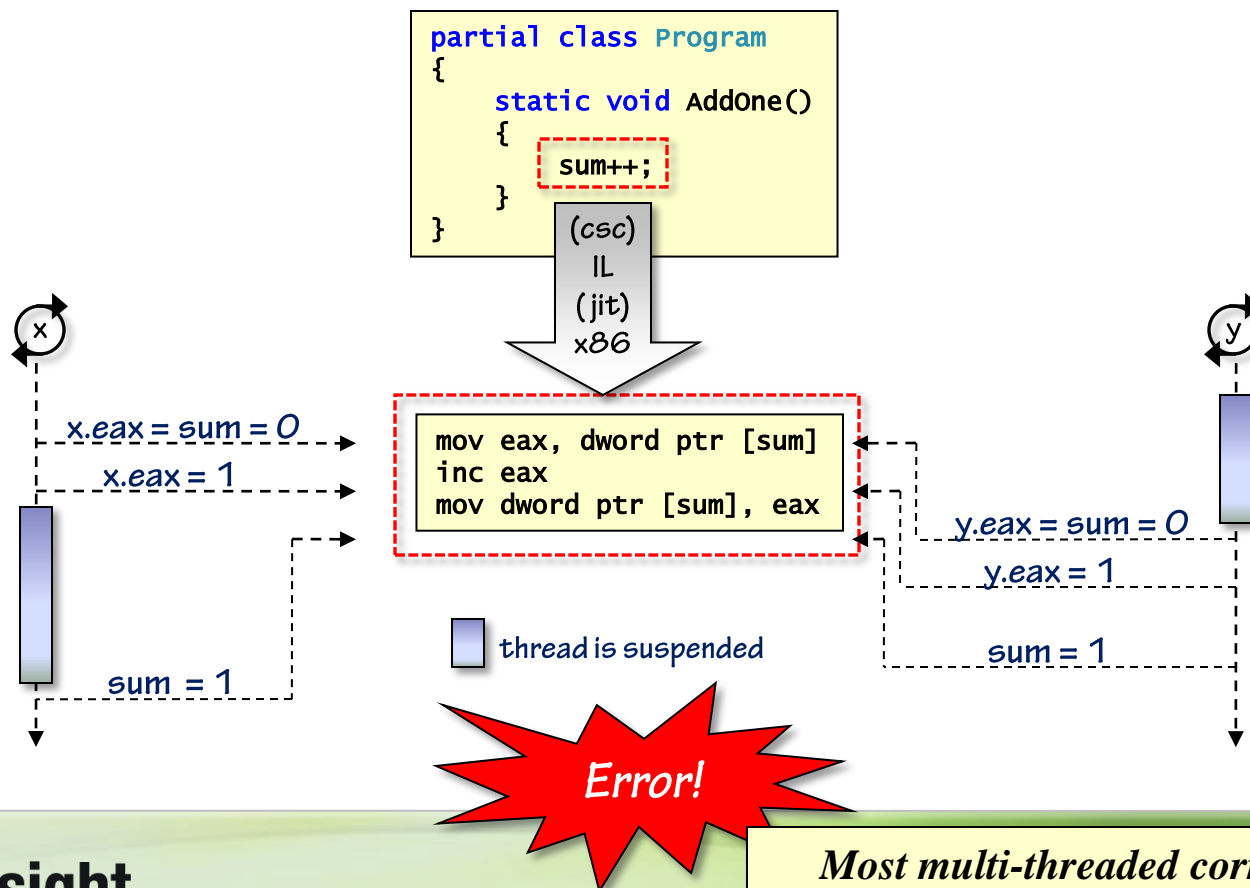
Critical Sections

- A *critical section* is a region of code that accesses a shared resource




Race Conditions

- A *race condition* occurs when the outcome may be affected by timing
 - e.g., uncontrolled access to a critical section



Solution #1: Atomic Updates

- Most processors support atomic updates of word-sized data



“lock” is an x86 instruction prefix that coordinates multi(core|processor) access to memory


- The FCL provides a processor-independent suite of atomic updates
 - ***System.Threading.Interlocked.Xxx***

```
partial class Program
{
    static void AddOne()
    {
        Interlocked.Increment(ref sum);
    }
}
```


*now a thread-safe
read/modify/write*

Solution #2: Data Partitioning

- Sometimes can partition data to orchestrate multi-threaded access
 - Depends on the data or resources being operated on
 - Requires problem-domain specific programming
 - “You work on this while I work on that” model
 - e.g., directory-based file operations



- e.g., array manipulations



Solution #3: Wait-Based Synchronization

- Sometimes threads need access to exact same resource — can't partition
 - e.g., insert or delete node from a list while other thread(s) are navigating list
 - e.g., multiple threads trying to manipulate the same file
- Sometimes data dependencies prevent a partitioned approach
 - When the output of one thread is required as input to another
 - e.g., computing the Fibonacci sequence
 - $\text{sequence}[0] = \text{sequence}[1] = 1$
 - $\text{sequence}[n] = \text{sequence}[n-1] + \text{sequence}[n-2]$
- Such situations require a wait-based approach to synchronization
 - i.e. thread may have to block until access is allowed...


Wait-Based Thread Synchronization

- Wait-based synchronization is based on a voluntary protocol

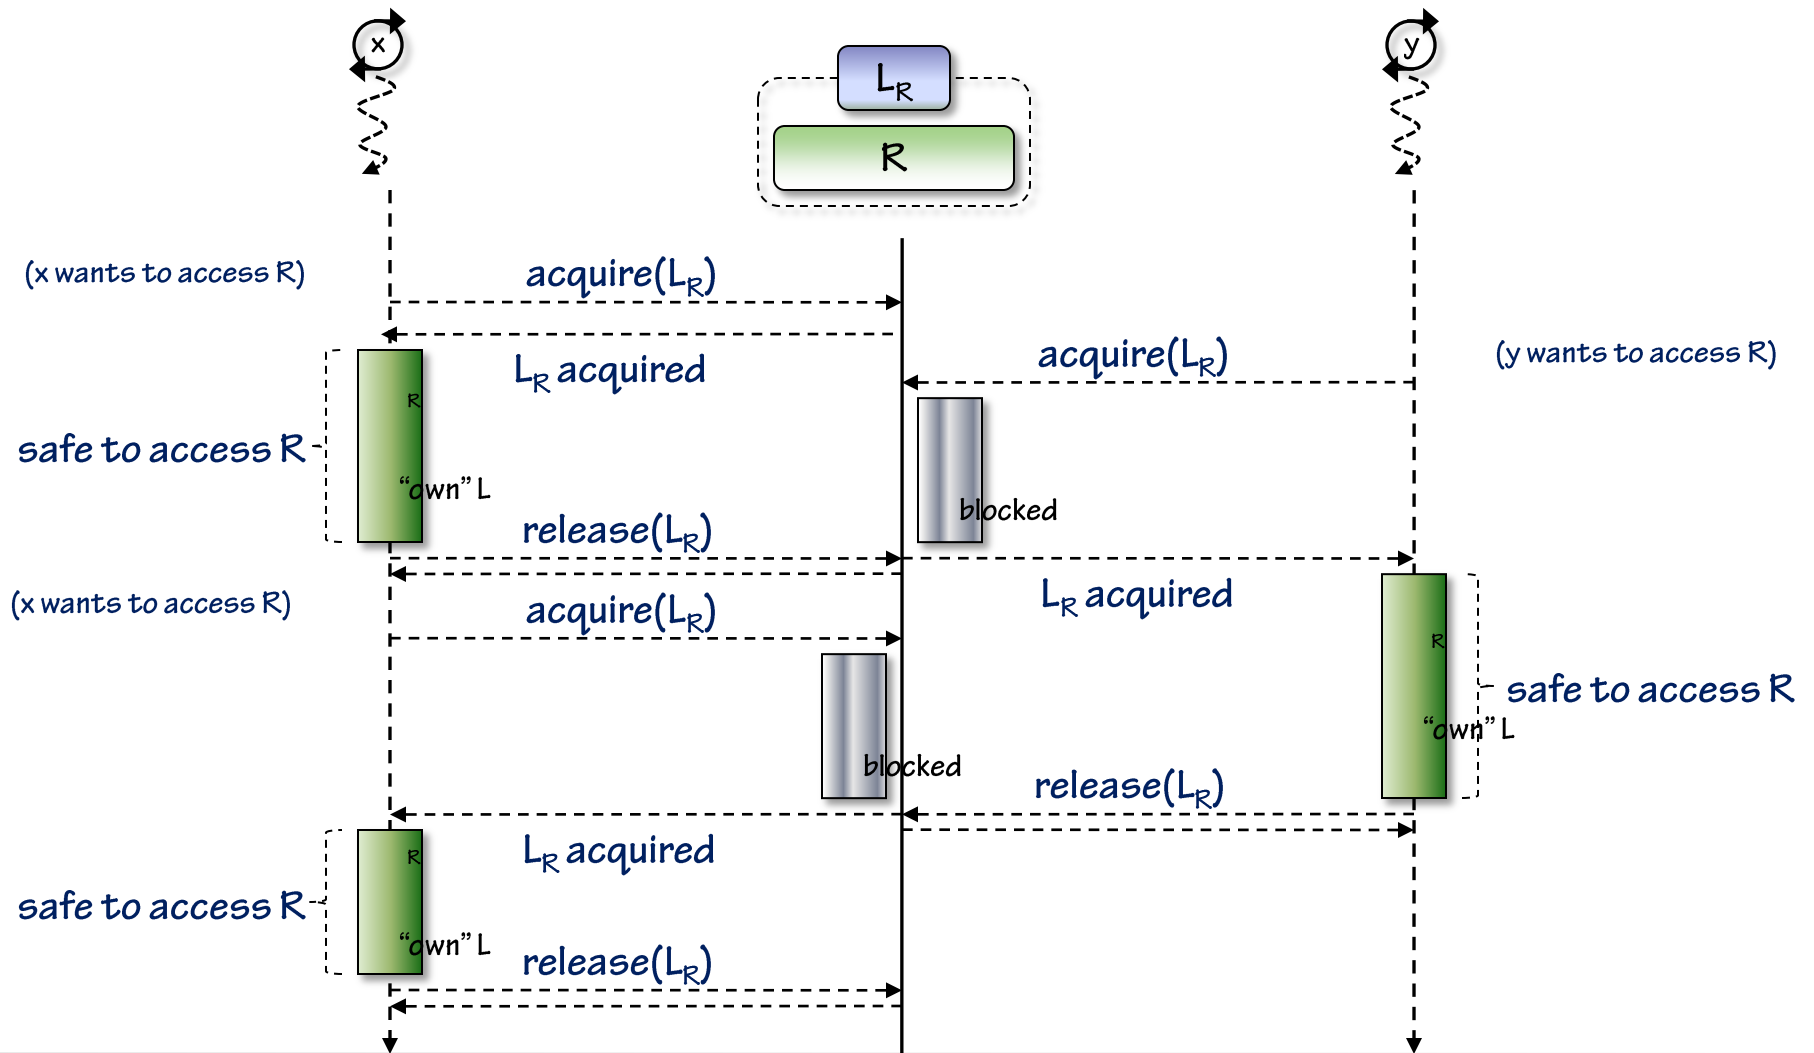
- A “gentleman’s agreement” or “handshake” model
- Elements of the protocol:
 - a shared resource is identified (“that array over there”)
 - a synchronization primitive/tool is agreed on (Monitor, mutex, ...)
 - an agreed upon *instance* of that primitive is identified
 - lock X guards file X, lock Y guards file Y, etc..
 - any thread wishing to accessing the resource agrees to:
 1. acquire ownership of the agreed upon synchronization primitive
 2. access the shared resource only after ownership acquired
 3. release ownership of the synchronization primitive once access is complete

a blocking
operation

- *Key point: the protocol is voluntary*



Wait-Based Synchronization



Wait-Based Synchronization Primitives

- Several wait-based synchronization primitives available in the CLR

- Monitor
- Mutex
- ReaderWriterLockSlim
- ManualResetEvent, AutoResetEvent
- Semaphore


*classes in
System.Threading
namespace*

- The first 3 share the same basic usage model

- Make a function call to acquire ownership of the “lock”
- Use the shared resource the designated “lock” is meant to protect
- Make a function call to release “lock” ownership once no longer needed

System.Threading.Monitor

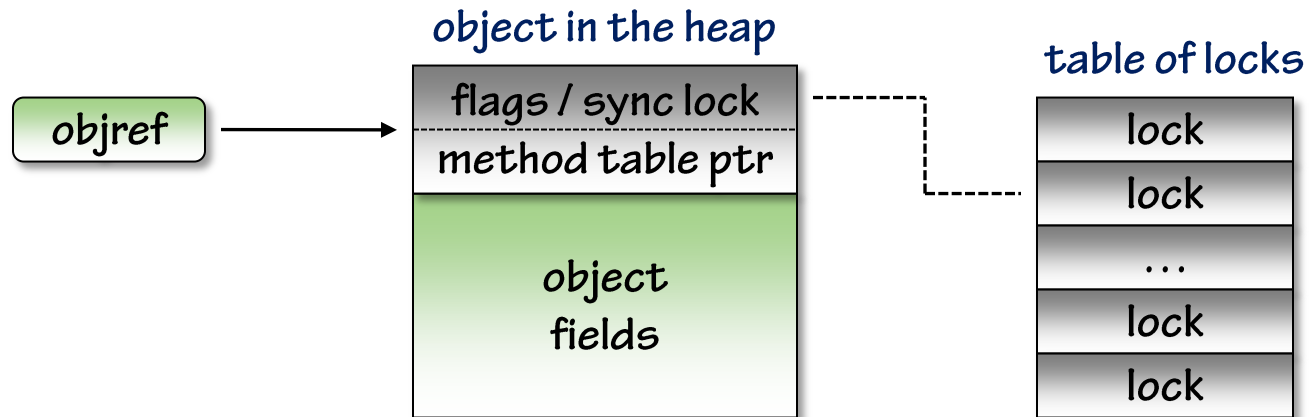
- **Monitors models gated access to a resource**
 - Threads agree to “enter” the monitor before accessing shared resource
 - CLR allows only one thread at a time to enter monitor
 - Other threads attempting to enter monitor while in use are blocked
 - May be recursively entered by same thread
 - Threads agree to “exit” the monitor once access to resource is complete
 - Next thread waiting for entry to monitor (if any) is allowed in
 - Recursive entrance operations by same thread require balanced exit operations



Monitors in the CLR

■ Monitor methods operate on object references

- Every object in the heap can *potentially* be associated with a lock
- Lock is demand-initialized by the CLR only if/when needed
- Index/reference to lock is stored within the CLR-managed object header
- Therefore, any object may effectively be used as a monitor



Implementation-specific details

Monitor Usage

```
public class widget2D
{
    int _x;
    int _y;

    public widget2D(int x, int y)
    {
        _x = x;
        _y = y;
    }

    public void MoveBy(int deltaX, int deltaY)
    {
        _x += deltaX;
        _y += deltaY;
    }

    public void GetPos(out int x, out int y)
    {
        x = _x;
        y = _y;
    }
}
```

critical sections

```
public class widget2D
{
    int _x;
    int _y;
    object _lock = new object();

    public widget2D(int x, int y)
    {
        _x = x;
        _y = y;
    }

    public void MoveBy(int deltaX, int deltaY)
    {
        Monitor.Enter(_lock);
        _x += deltaX;
        _y += deltaY;
        Monitor.Exit(_lock);
    }

    public void GetPos(out int x, out int y)
    {
        Monitor.Enter(_lock);
        x = _x;
        y = _y;
        Monitor.Exit(_lock);
    }
}
```

used to gate access to widget fields

access to fields is synchronized

access to fields is synchronized

Exception-Safe Monitor Usage

```
public class widget2D
{
    int _x;
    int _y;
    object _lock = new object();

    public widget2D(int x, int y)
    {
        _x = x;
        _y = y;
    }

    public void MoveBy(int deltaX, int deltaY)
    {
        Monitor.Enter(_lock);
        _x += deltaX;
        _y += deltaY;
        Monitor.Exit(_lock);
    }

    public void GetPos(out int x, out int y)
    {
        Monitor.Enter(_lock);
        x = _x;
        y = _y;
        Monitor.Exit(_lock);
    }
}
```

exception-prone

exception-safe

```
public class widget2D
{
    int _x;
    int _y;
    object _lock = new object();

    public void MoveBy(int deltaX, int deltaY)
    {
        Monitor.Enter(_lock);

        try {
            _x += deltaX;
            _y += deltaY;
        }
        finally {
            Monitor.Exit(_lock);
        }
    }

    public void GetPos(out int x, out int y)
    {
        Monitor.Enter(_lock);

        try {
            x = _x;
            y = _y;
        }
        finally {
            Monitor.Exit(_lock);
        }
    }
}
```

C# Monitor Usage

```
public class widget2D
{
    int _x;
    int _y;
    object _lock = new object();

    public void MoveBy(int deltaX, int deltaY)
    {
        Monitor.Enter(_lock);

        try {
            _x += deltaX;
            _y += deltaY;
        }
        finally {
            Monitor.Exit(_lock);
        }
    }

    public void GetPos(out int x, out int y)
    {
        Monitor.Enter(_lock);

        try {
            x = _x;
            y = _y;
        }
        finally {
            Monitor.Exit(_lock);
        }
    }
}
```

C# equivalent

```
public class widget2D
{
    int _x;
    int _y;
    object _lock = new object();

    public void MoveBy(int deltaX, int deltaY)
    {
        lock(_lock)
        {
            _x += deltaX;
            _y += deltaY;
        }
    }

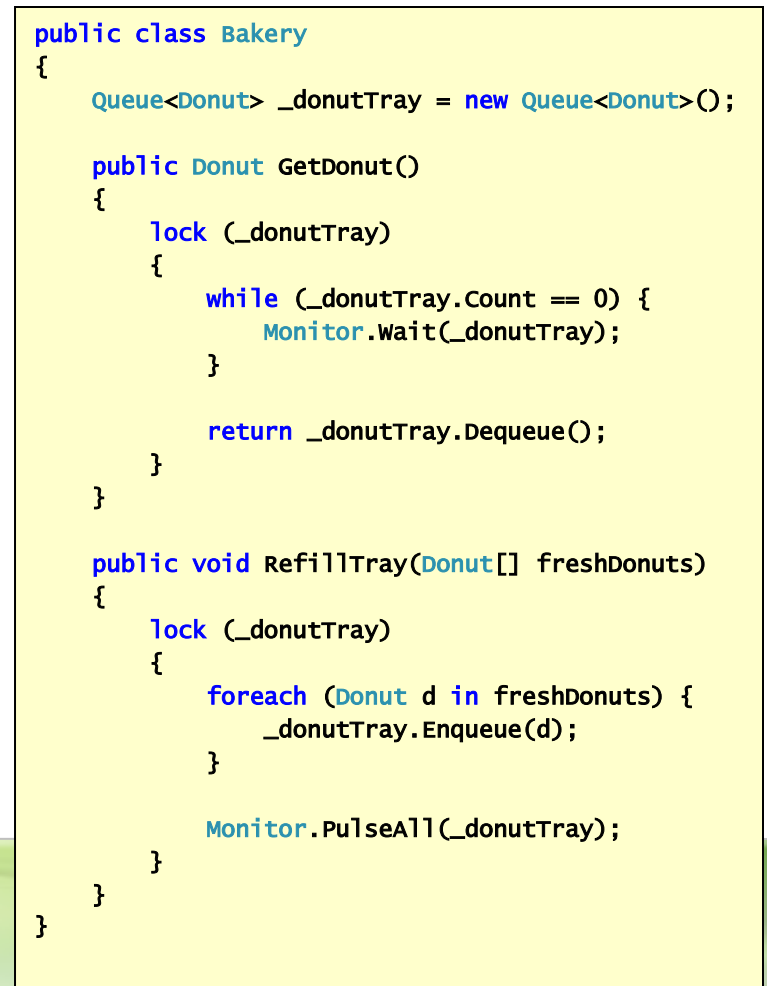
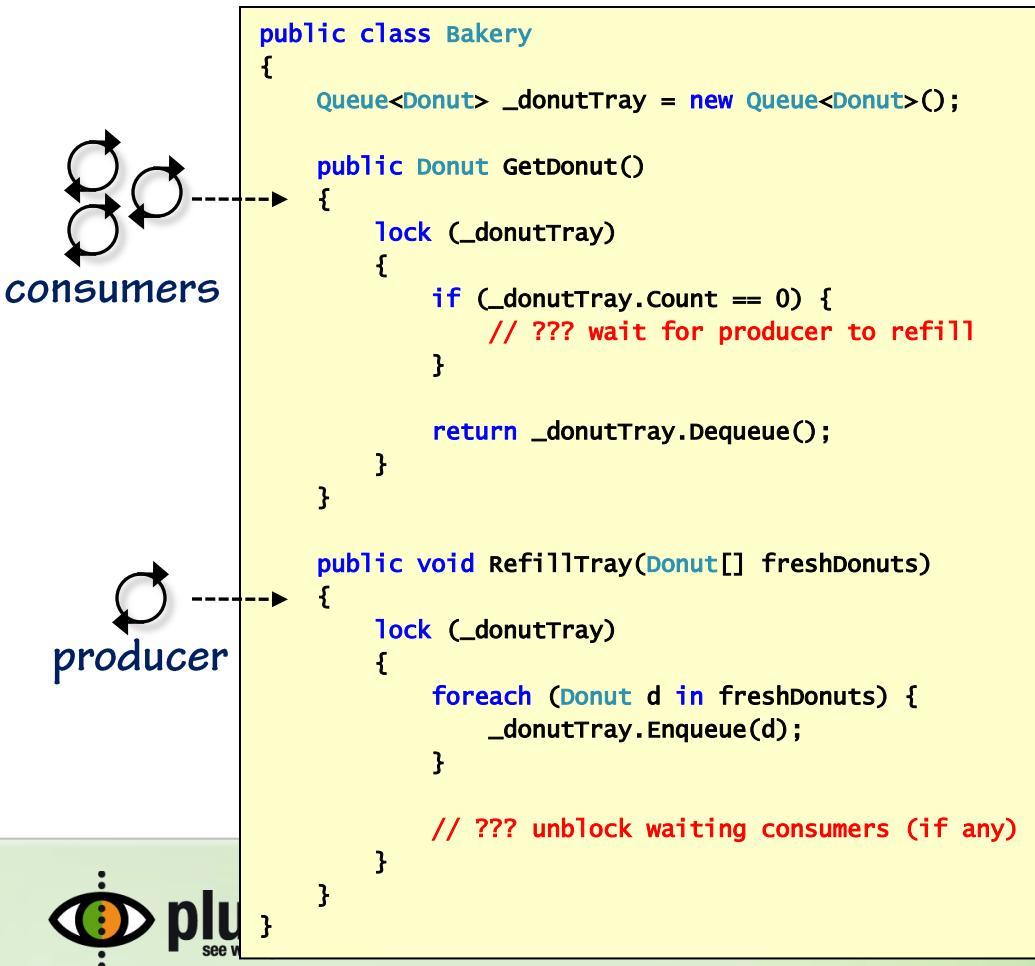
    public void GetPos(out int x, out int y)
    {
        lock(_lock)
        {
            x = _x;
            y = _y;
        }
    }
}
```


Hold & Wait

- **Sometimes a thread needs to wait for something *while* holding a lock**
 - e.g.: a resource to be provided or replenished
 - e.g.: another lock
- **Resource replenishment or other condition**
 - producer/consumer model for resource handling
 - blocking semantics for consumers when resource(s) not available
- **Multiple lock acquisition**
 - atomic updates of multiple resources, each protected by own lock

Hold & Wait with Monitors

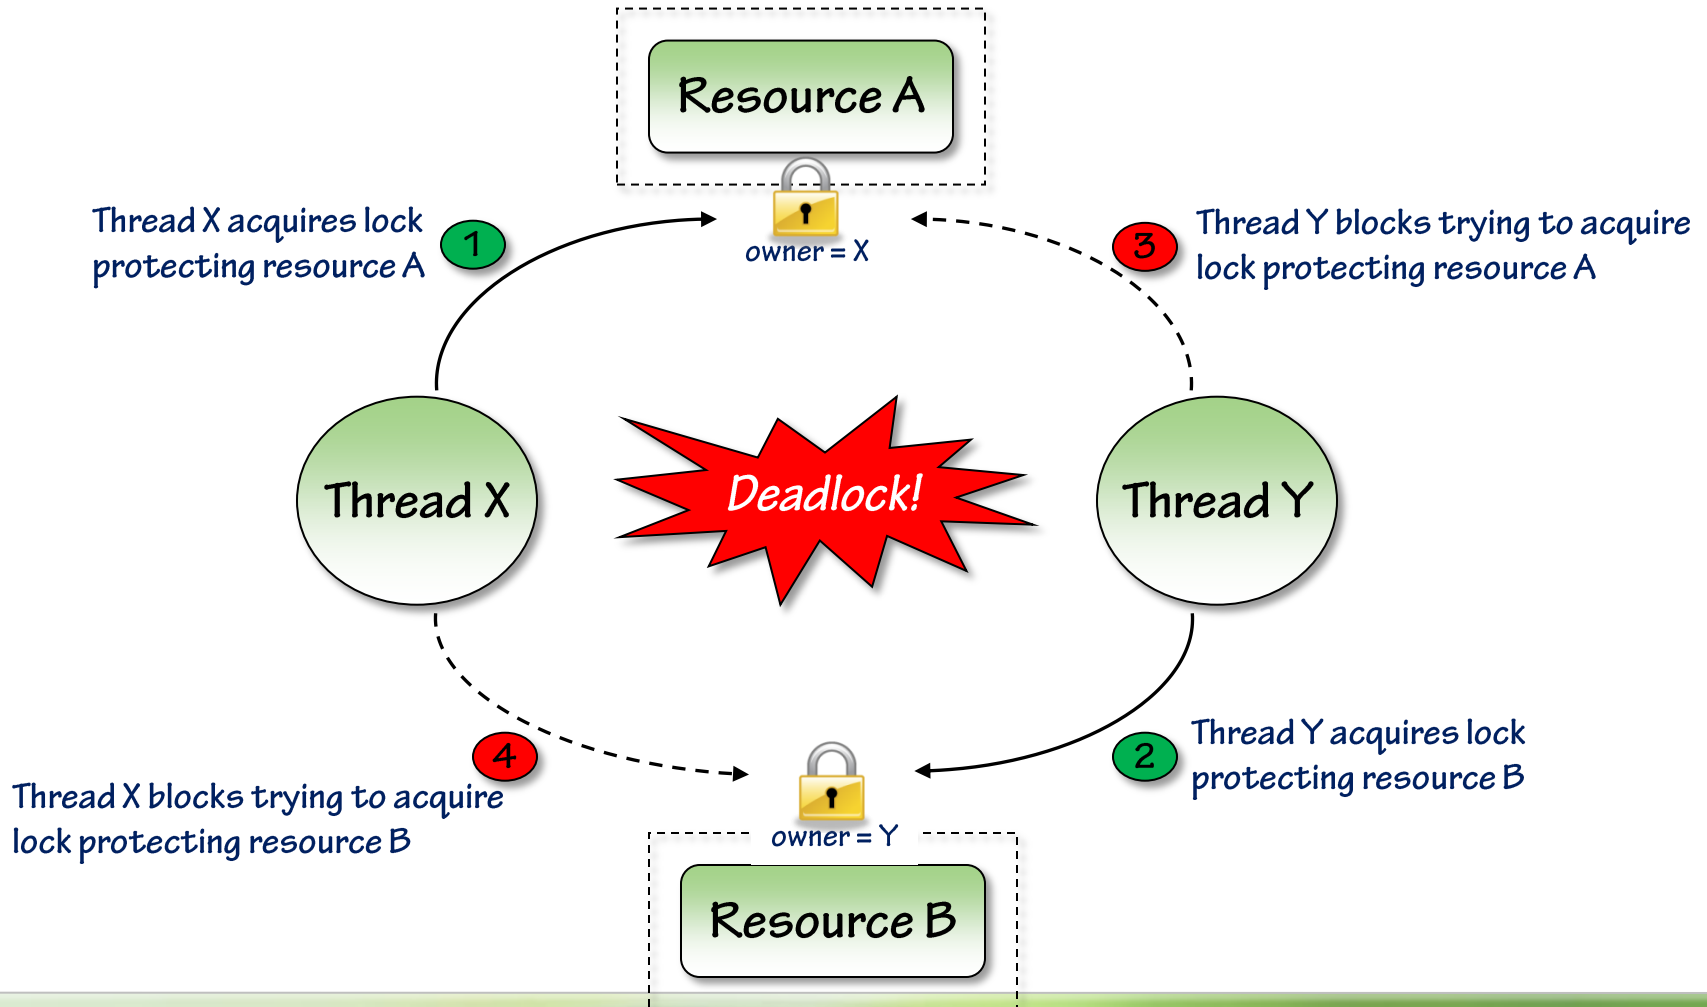
- The Monitor class supports safe hold-and-wait operations
 - Via the Pulse, PulseAll, Wait methods



Deadlock

- **Deadlock can occur whenever a *hold-and-wait* situation is possible**
 - While holding one lock, a thread attempts to acquire another lock
- **Note that the above is rife with caveats:**
 - Deadlock can occur (but won't necessarily occur)
 - requires 2 or more threads competing for the same set of locks
 - Deadlock is possible (but not necessarily probable)
 - probability increases as # of threads/processors/cores increases
 - Deadlock may only be temporary
 - if timeouts are used in all lock acquisition calls

Deadlock Illustrated



Mutexes

- **A mutex is a Win32 kernel object**
 - ***System.Threading.Mutex*** provides an FCL wrapper for managed code
- **Benefits**
 - Supports timeout-limited lock acquisition
 - Nameable; enabling cross-process (same-machine) thread synchronization
 - Enables deadlock-free multiple lock acquisition via ***WaitHandle.WaitAll***
- **Tradeoffs**
 - Acquisition & release calls always incur roundtrip to/from kernel mode
 - Underlying kernel object must be closed when no longer needed
 - handled automatically (but on a delayed basis) by GC/finalization mechanics

Deadlock-Prone Multiple Lock Acquisition

```
public class BankAccount
{
    double _balance;
    object _lock = new object();

    public void Credit(double amt)
    {
        lock (_lock)
        {
            _balance += amt;
        }
    }

    public void Debit(double amt)
    {
        lock (_lock)
        {
            _balance -= amt;
        }
    }

    public void TransferFrom(BankAccount otherAcct, double amt)
    {
        lock (this._lock)
        {
            lock (otherAcct._lock)
            {
                otherAcct.Debit(amt);
                this.Credit(amt);
            }
        }
    }
}
```

thread-safe

*Vulnerable to
deadlock!*

Hold-and-wait
potential exists
based on timing

Deadlock-Free Multiple Lock Acquisition

```
public partial class BankAccount
```

```
{  
    double _balance;  
    Mutex _lock = new Mutex();
```

now using a Mutex

```
    public void Credit(double amt)  
    {
```

```
        if (_lock.WaitOne())  
        {
```

```
            try  
            {  
                _balance += amt;
```

```
            }  
            finally
```

```
            {  
                _lock.ReleaseMutex();  
            }  
        }  
    }
```

```
    public void Debit(double amt)  
    {
```

```
        if (_lock.WaitOne())  
        {
```

```
            try  
            {  
                _balance -= amt;
```

```
            }  
            finally
```

```
            {  
                _lock.ReleaseMutex();  
            }  
        }  
    }
```

based lock acquisition

Mutex

```
public partial class BankAccount
```

```
{  
    public void TransferFrom(BankAccount otherAcct, double amt)  
    {
```

```
        Mutex[] locks = { this._lock, otherAcct._lock };
```

```
        if (WaitHandle.WaitAll(locks))
```

```
        {
```

```
            try
```

```
            {  
                otherAcct.Debit(amt);  
                this.Credit(amt);
```

```
            }  
            finally
```

```
            {
```

```
                foreach (Mutex m in locks)
```

```
                {
```

```
                    m.ReleaseMutex();
```

```
                }  
            }  
        }
```

```
    }
```

all locks
acquired
without risk
of deadlock



Summary

- **Most resources are not meant to be accessed concurrently**
 - Improper access yields race conditions and incorrect results
 - very difficult to detect, debug, and fix
- **Common solutions**
 - Atomic updates of machine-word-sized values
 - Data partitioning
 - Wait-based synchronization (Monitor, Mutex, others)
- **Beware**
 - Critical sections
 - Race conditions
 - Deadlock