

# xml indexes and xquery performance

Bob Beauchemin



# Overview

- **XQuery methods and processing**
  - Querying non-indexed XML
- **XML Indexes**
  - Kinds of XML Indexes
  - Creating and Maintaining
  - The XML Node Table
  - Querying indexed XML
- **Usage of XML Indexes**
  - Query plan changes
  - Which index for which use case
- **XQuery Performance Hints**

# XQuery and SQL

- **XQuery uses the relational query engine**
  - XQuery methods are SQL "system defined functions"
  - Integrated query plan is produced
  - XQuery operations added to (relational) query plan
    - Shows up in query plans as UDX
  - Types of UDX operators
    - Check UDX - validates XML being inserted
    - XQuery Serializer UDX - serializes the query result as XML
    - XQuery Data UDX - evaluates the XQuery data() function
    - XQuery String UDX - evaluates the XQuery string() function
    - XQuery Contains UDX - evaluates the XQuery contains() function
    - FOR XML UDX – used in SELECT... FOR XML construction

# XQuery Methods and Processing

- **XQuery methods are part of a specific SQL statement**
  - Both SQL and XQuery go through a static phase
  - SQL static phase includes
    - SQL Parser
    - Static Typing - using SQL metadata
    - Algebrization
  - XQuery static phase includes
    - XQuery Parser
    - Static Typing - using (optional) XML Schema Collection
    - Algebrization
  - Trees are combined optimized
    - Static optimization of the combined logical and physical tree
    - Runtime optimization and execution of the tree - using relational & XML indexes

# The Node Table

- **To be useable by SQL, XML is decomposed into a “node table”**
  - Node table is a relational table
  - Visible in sys.internal\_tables if materialized in primary XML index
  - Almost always used in XQuery
  - One row per node - possible overflow rows
- **Appears as "XmlReader with XPath Filter"**
  - Similar but not the same as System.Xml.XmlReader
- **Non-indexed XML – no node cardinality estimates**
  - Decomposition at execution time

# XML Query Plans

- Query plans against table with XML column can be categorized
- This query can result in either top-down plan
  - Select rows in base table that qualify
  - Do XQuery against XML instance in qualifying
- Or bottom-up plan
  - Do XQuery first
  - Then join to base table

```
-- get qualifying (non-NULL) rows first?  
-- or get InvoiceID 1003 first?  
SELECT xml_col.exist('/Invoice/@InvoiceID[.= "1003"]')  
FROM xml_tab  
GO
```

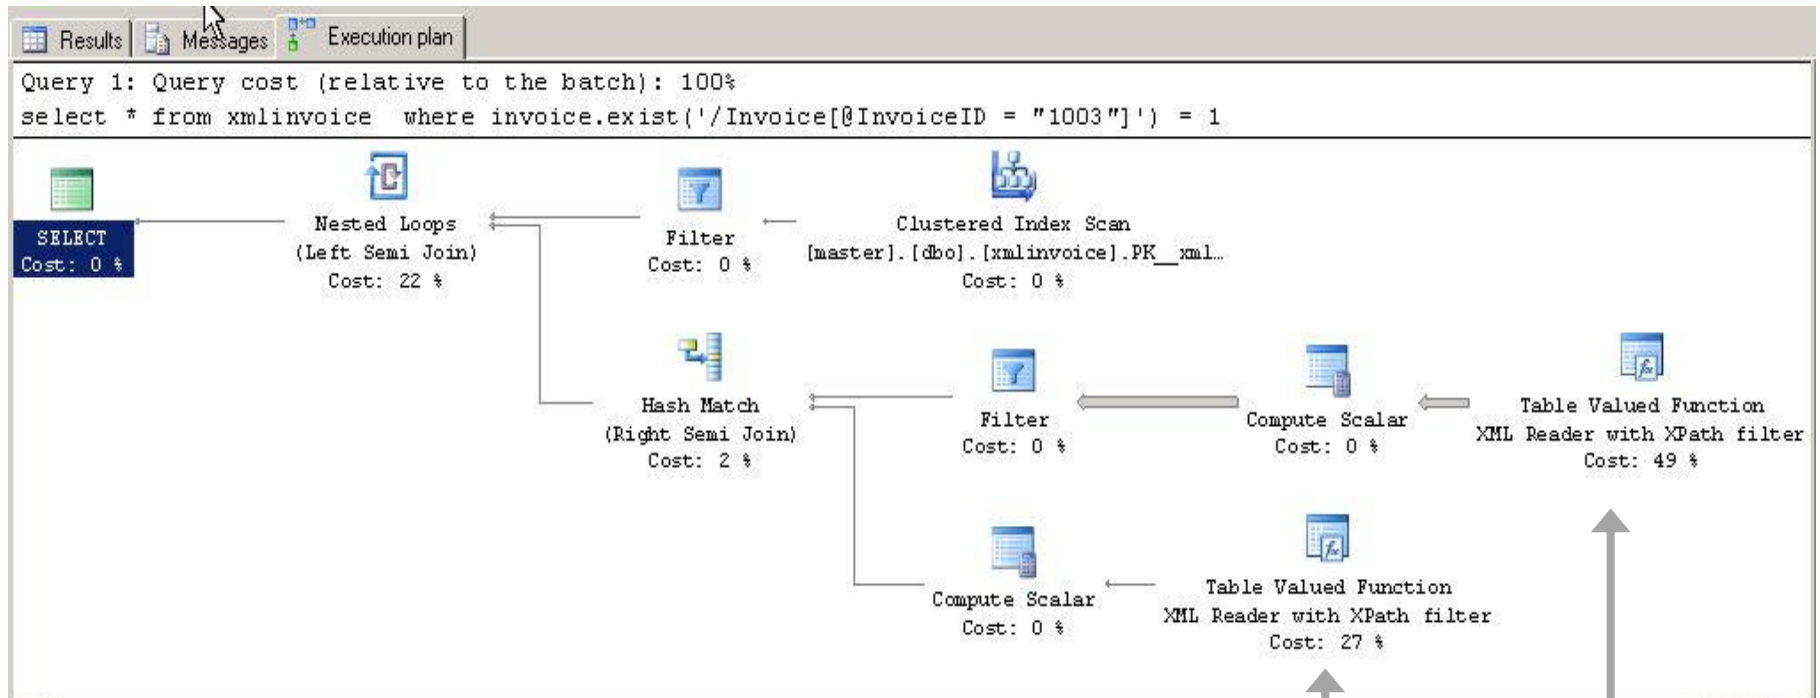
# Non-indexed XML and Evaluation Steps

- **Equivalent evaluations are converted to relational plans differently**
  - prefer predicates that use current node over last location step of path
  - This only works when the node returned is irrelevant (when using xml.exist)

```
-- this syntax preferred
SELECT xml_col.exist('/Invoice/@InvoiceID[.= "1003"]')
FROM xml_tab
GO
```

```
-- over this query
-- note: this does not produce the same result
-- except when using xml.exist
SELECT xml_col.exist('/Invoice[@InvoiceID = "1003"]')
FROM xml_tab
GO
```

# Query Plan Without Index



Multiple Evaluation Steps  
Using XML Reader



# XML Indexes

- **XML INDEXes are special indexes on an XML column**
  - Optimizes XQuery operations on the column
  - Primary XML index must be created first
    - Before any additional XML indexes can be created
  - Table must have clustered index
    - XML column cannot be the clustered index
  - Primary XML index provides cardinality estimates for the query optimizer

```
CREATE TABLE xml_tab (  
    id integer primary key,  
    doc xml)  
GO  
CREATE PRIMARY XML INDEX xml_idx on xml_tab (doc)  
GO
```

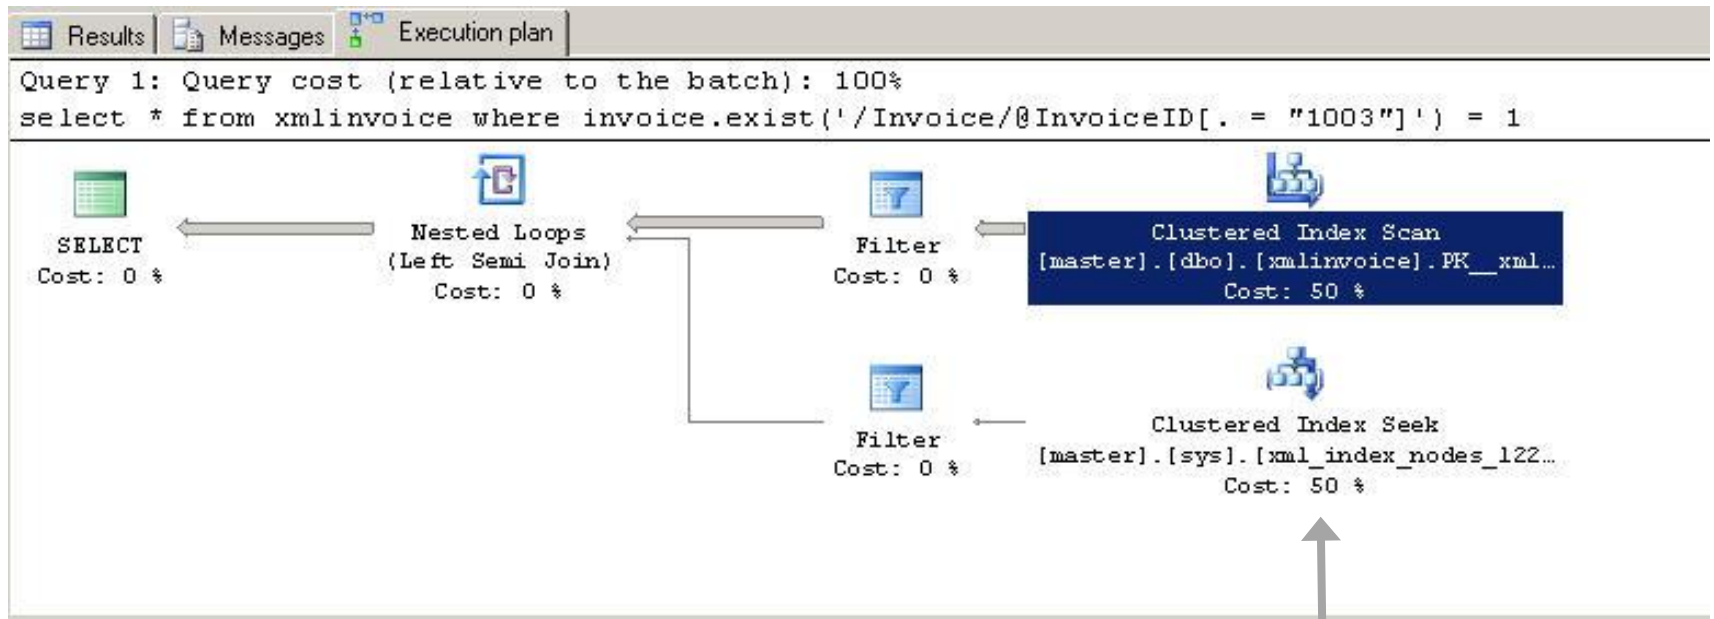
# XML Primary Index Details

- **Node table is materialized**
  - Base table clustered index, plus 12 columns of info
  - Column metadata visible with join of sys.indexes and sys.columns
  - Clustered index on
    - primary key (of base table)
    - id (node\_id of the node table)
- **Most of the same create/alter options as SQL index**
- **Only available on columns (not variables)**
- **Space utilization is 2-5 times original data**
  - Can slow down inserts

# Columns in the Node Table

Column Name	Column Description	Data Type
id	node identifier in ordpath format	varbinary(900)
nid	node name (tokenized)	int
tagname	tag name	nvarchar(4000)
taguri	tag uri	nvarchar(4000)
tid	node data type (tokenized)	int
value	first portion of the node value	sql_variant
lvalue	long node value (pointer)	nvarchar(max)
lvaluebin	long node value in binary (pointer)	varbinary(max)
hid	path (tokenized)	varchar(900)
xsinil	is it NULL (xsi:nil)	bit
xsitype	does it use xsi:type	bit
pk1	primary key of the base table	int

# Using Primary XML Index



Clustered Index Seek  
On Primary XML Index

# Secondary XML Indexes

- Three specialized index types
  - VALUE – optimizes content queries
  - PATH – optimizes when path selective
  - PROPERTY – top-down queries, i.e. rows selected first

```
CREATE TABLE xml_tab (  
    id integer primary key,  
    doc xml)  
GO  
CREATE PRIMARY XML INDEX xml_idx on xml_tab (doc)  
GO  
-- secondary index  
CREATE XML INDEX path_idx on xml_tab (doc)  
    USING XML INDEX xml_idx FOR PATH  
GO
```

# Secondary XML Indexes

- **Secondaries are non-clustered indexes on the node table**
  - PATH is index on (HID, value)
    - HID is the path
  - VALUE is index on (value, HID)
    - Same columns as PATH, but reverse order
  - PROPERTY is index on (pk, HID, value)
    - Same as path but includes base primary key

# XML Index Maintenance

- **XML Indexes are efficiently maintained on update**
  - Key part - maintained like relational index
  - Path part - efficient because of ORDPATH representation
    - Only the changes part of the document is built
    - Semantics similar to XML modify
- **On insert, XML must be decomposed to build new index rows**

# Which Index?

- **If all indexes are present**

- Primary index is "default" index
  - Used if no other index selective enough
  - Used for construction
    - with special value 'DESC'
- VALUE is often used with wildcard queries
  - /SomeNode/@\*[. = "special"]
  - (i.e. any attribute with the value special)
- PATH used only if ragged documents
  - many different paths, high path selectivity
- PROPERTY used with sparse attributes
  - when rows are selected first
  - when '/' appears in the path
  - many different name-value pairs, different names



# Path and Axis Performance Hints

- **Queries that use parent axis generate extra query plan steps**
  - Avoid parent axis queries if possible
  - Use multiple CROSS APPLY steps rather than parent axis to get nodes at multiple nesting levels using xml.nodes
- **Move ordinals to the end of path expressions**
  - Ensure that the answer is equivalent
- **Avoid predicates in the middle of path expressions**
- **Specify a single root node in query**
  - Optimizer assumes XML can be a fragment

# More XQuery Performance Hints

- **Use xml.value rather than xml.query and a cast**
  - Almost always produce same answer
- **Casting in XQuery is expensive**
  - Prohibits index use
- **Avoid functions that aggregate all text nodes if there is a single text node**
  - Prohibits index use
- **Avoid multiple XQuery evaluations in a single SQL query**
  - Use SQL subqueries instead
- **Avoid applying inline functions to XQuery results**
- **When joining SQL and XML values,**
  - xml.exist and sql:column - better
  - xml.value and comparison to SQL column - worse

# Constructing XML from SQL Values

- **XML Path can be quicker for construction**
  - Composing XML using XQuery most useful for transforming existing document, not construction from SQL values
  - Composing XML nodes from multiple relational columns better using FOR XML PATH
  - Composing XML using XML DML insert node-at-a-time is slowest

# Optimizations in xml.modify

- **The xml.modify method uses**
  - sparse updates
  - sparse logging
- **Concurrency is still at row (document) level**
- **For best results**
  - insert new nodes as last sibling
  - insert/update nodes deeper in the hierarchy

# Review

- **XML Indexes speed up XQuery**
  - XQuery uses relational query engine
- **Primary XML Index creates node table**
  - Primary XML Index refers to clustered index on node table
  - Secondary are non-clustered index
- **Primary XML Index almost always useful**
  - 2-5 times size of document
- **Secondary XML Indexes use-case based**
  - Like SQL indexes
- **Phrasing of XQuery can effect performance**

# References

- [XML Indexes in SQL Server 2005](#), Bob Beauchemin, MSDN Online
- [XQuery Implementation in a Relational Database System](#), Shankar Pal, et al., Conference on VLDB
- [Performance Optimizations for the XML Data Type in SQL Server](#), Shankar Pal, et al., MSDN Online
- SQL Server 2005 XQuery Performance Tips, Ward Pond
  - <http://blogs.technet.com/b/wardpond/archive/2005/06/23/sql-server-2005-xquery-performance-tips.aspx>
- MSDN Webcast: Making the Most of XQuery with SQL Server 2005, Michael Rys

For more in-depth **online** developer **training** visit



**on-demand** content from authors you **trust**